



Scipy.org (<https://scipy.org/>)    Docs (<https://docs.scipy.org/>)

SciPy v0.19.0 Reference Guide ([../index.html](https://docs.scipy.org/doc/scipy-0.19.0/index.html))

Optimization and root finding (**scipy.optimize**) ([../optimize.html](https://docs.scipy.org/doc/scipy-0.19.0/optimize/index.html))

index ([../genindex.html](https://docs.scipy.org/doc/scipy-0.19.0/optimize/genindex.html))    modules ([../py-modindex.html](https://docs.scipy.org/doc/scipy-0.19.0/optimize/py-modindex.html))

modules ([../scipy-optimize-modindex.html](https://docs.scipy.org/doc/scipy-0.19.0/optimize/scipy-optimize-modindex.html))    next ([scipy.optimize.minimize\\_scalar.html](https://docs.scipy.org/doc/scipy-0.19.0/optimize/scipy.optimize.minimize_scalar.html))

previous ([../optimize.nonlin.html](https://docs.scipy.org/doc/scipy-0.19.0/optimize/scipy.optimize.optimize.nonlin.html))

## scipy.optimize.minimize

**scipy.optimize.minimize**(*fun*, *x0*, *args=()*, *method=None*,  
*jac=None*, *hess=None*, *hessp=None*, *bounds=None*, *constraints=()*,  
*tol=None*, *callback=None*, *options=None*)

[[source](#)]

([http://github.com/scipy/scipy/blob/v0.19.0/scipy/optimize/\\_minimize.py#L36-L466](http://github.com/scipy/scipy/blob/v0.19.0/scipy/optimize/_minimize.py#L36-L466))

Previous topic

Nonlinear solvers  
([../optimize.nonlin.ht](https://docs.scipy.org/doc/scipy-0.19.0/optimize/optimize.nonlin.html))

Next topic

[scipy.optimize.minir](#)  
([scipy.optimize.minir](https://docs.scipy.org/doc/scipy-0.19.0/optimize/scipy.optimize.minir.html))

Minimization of scalar function of one or more variables.

In general, the optimization problems are of the form:

minimize  $f(x)$  subject to

$g_i(x) \geq 0, \quad i = 1, \dots, m$

$h_j(x) = 0, \quad j = 1, \dots, p$

where  $x$  is a vector of one or more variables.  $g_i(x)$  are the inequality constraints.  $h_j(x)$  are the equality constraints.

Optionally, the lower and upper bounds for each element in  $x$  can also be specified using the *bounds* argument.

**Parameters:** *fun* : *callable*

Objective function.

*x0* : *ndarray*

Initial guess.

*args* : *tuple, optional*

Extra arguments passed to the objective function and its derivatives (Jacobian, Hessian).

*method* : *str or callable, optional*

Type of solver. Should be one of

- 'Nelder-Mead' (*see here*) (`../optimize.minimize-neldermead.html#optimize-minimize-neldermead`)
- 'Powell' (*see here*) (`../optimize.minimize-powell.html#optimize-minimize-powell`)
- 'CG' (*see here*) (`../optimize.minimize-cg.html#optimize-minimize-cg`)
- 'BFGS' (*see here*) (`../optimize.minimize-bfgs.html#optimize-minimize-bfgs`)
- 'Newton-CG' (*see here*) (`../optimize.minimize-newtoncg.html#optimize-minimize-newtoncg`)
- 'L-BFGS-B' (*see here*) (`../optimize.minimize-lbfgsb.html#optimize-minimize-lbfgsb`)
- 'TNC' (*see here*) (`../optimize.minimize-tnc.html#optimize-minimize-tnc`)
- 'COBYLA' (*see here*) (`../optimize.minimize-cobyla.html#optimize-minimize-cobyla`)
- 'SLSQP' (*see here*) (`../optimize.minimize-slsqp.html#optimize-minimize-slsqp`)
- 'dogleg' (*see here*) (`../optimize.minimize-dogleg.html#optimize-minimize-dogleg`)
- 'trust-ncg' (*see here*) (`../optimize.minimize-trustncg.html#optimize-minimize-trustncg`)
- custom - a callable object (added in version 0.14.0), see below for description.

If not given, chosen to be one of BFGS, L-BFGS-B, SLSQP, depending if the problem has constraints or bounds.

***jac*** : *bool or callable, optional*

Jacobian (gradient) of objective function. Only for CG, BFGS, Newton-CG, L-BFGS-B, TNC, SLSQP, dogleg, trust-ncg. If *jac* is a Boolean and is True, *fun* is assumed to return the gradient along with the objective function. If False, the gradient will be estimated numerically. *jac* can also be a callable returning the gradient of the objective. In this case, it must accept the same arguments as *fun*.

***hess, hessp*** : *callable, optional*

Hessian (matrix of second-order derivatives) of objective function or Hessian of objective function times an arbitrary vector *p*. Only for Newton-CG, dogleg, trust-ncg. Only one of *hessp* or *hess* needs to be given. If *hess* is provided, then *hessp* will be ignored. If neither *hess* nor *hessp* is provided, then the Hessian product will be approximated using finite differences on *jac*. *hessp* must compute the Hessian times an arbitrary vector.

***bounds*** : *sequence, optional*

Bounds for variables (only for L-BFGS-B, TNC and SLSQP). (*min*, *max*) pairs for each element in *x*, defining the bounds on that parameter. Use None for one of *min* or *max* when there is no bound in that direction.

**constraints** : *dict or sequence of dict, optional*

Constraints definition (only for COBYLA and SLSQP). Each constraint is defined in a dictionary with fields:

**type** : *str*

Constraint type: 'eq' for equality, 'ineq' for inequality.

**fun** : *callable*

The function defining the constraint.

**jac** : *callable, optional*

The Jacobian of *fun* (only for SLSQP).

**args** : *sequence, optional*

Extra arguments to be passed to the function and Jacobian.

Equality constraint means that the constraint function result is to be zero whereas inequality means that it is to be non-negative. Note that COBYLA only supports inequality constraints.

**tol** : *float, optional*

Tolerance for termination. For detailed control, use solver-specific options.

**options** : *dict, optional*

A dictionary of solver options. All methods accept the following generic options:

**maxiter** : *int*

Maximum number of iterations to perform.

**disp** : *bool*

Set to True to print convergence messages.

For method-specific options, see `show_options` (`scipy.optimize.show_options.html#scipy.optimize.show_options`).

**callback** : *callable, optional*

Called after each iteration, as `callback(xk)`, where `xk` is the current parameter vector.

**Returns:**

**res** : *OptimizeResult*

The optimization result represented as a `OptimizeResult` object. Important attributes are: `x` the solution array, `success` a Boolean flag indicating if the optimizer exited successfully and `message` which describes the cause of the termination. See `OptimizeResult` (`scipy.optimize.OptimizeResult.html#scipy.optimize.OptimizeResult`) for a description of other attributes.

**See also:**

`minimize_scalar`

([scipy.optimize.minimize\\_scalar.html#scipy.optimize.minimize\\_scalar](#))

Interface to minimization algorithms for scalar univariate functions

`show_options`

([scipy.optimize.show\\_options.html#scipy.optimize.show\\_options](#))

Additional options accepted by the solvers

## Notes

---

This section describes the available solvers that can be selected by the 'method' parameter. The default method is *BFGS*.

### Unconstrained minimization

Method *Nelder-Mead* ([../optimize.minimize-neldermead.html#optimize-minimize-neldermead](#)) uses the Simplex algorithm [R174], [R175]. This algorithm is robust in many applications. However, if numerical computation of derivative can be trusted, other algorithms using the first and/or second derivatives information might be preferred for their better performance in general.

Method *Powell* ([../optimize.minimize-powell.html#optimize-minimize-powell](#)) is a modification of Powell's method [R176], [R177] which is a conjugate direction method. It performs sequential one-dimensional minimizations along each vector of the directions set (*direc* field in *options* and *info*), which is updated at each iteration of the main minimization loop. The function need not be differentiable, and **no derivatives are taken**.

Method *CG* ([../optimize.minimize-cg.html#optimize-minimize-cg](#)) uses a nonlinear conjugate gradient algorithm by Polak and Ribiere, a variant of the Fletcher-Reeves method described in [R178] pp. 120-122. Only the first derivatives are used.

Method *BFGS* ([../optimize.minimize-bfgs.html#optimize-minimize-bfgs](#)) uses the quasi-Newton method of Broyden, Fletcher, Goldfarb, and Shanno (BFGS) [R178] pp. 136. It uses the first derivatives only. BFGS has proven good performance even for non-smooth optimizations. This method also returns an approximation of the Hessian inverse, stored as *hess\_inv* in the *OptimizeResult* object.

Method *Newton-CG* ([../optimize.minimize-newtoncg.html#optimize-minimize-newtoncg](#)) uses a Newton-CG algorithm [R178] pp. 168 (also known as the truncated Newton method). It uses a CG method to compute the search direction. See also *TNC* method for a box-constrained minimization with a similar algorithm.

Method *dogleg* ([../optimize.minimize-dogleg.html#optimize-minimize-dogleg](#)) uses the dog-leg trust-region algorithm [R178] for unconstrained minimization. This algorithm requires the gradient and Hessian; furthermore the Hessian is required to be **positive definite**.

Method *trust-ncg* ([../optimize.minimize-trustncg.html#optimize-minimize-trustncg](#)) uses the Newton conjugate gradient trust-region algorithm [R178] for unconstrained minimization. This algorithm requires the gradient and either the Hessian or a function that computes the product of the Hessian with a given vector.

### Constrained minimization

Method *L-BFGS-B* ([../optimize.minimize-lbfgsb.html#optimize-minimize-lbfgsb](#)) uses the L-BFGS-B algorithm [R179], [R180] for bound constrained minimization.

Method *TNC* ([../optimize.minimize-tnc.html#optimize-minimize-tnc](#)) uses a truncated Newton algorithm [R178], [R181] to minimize a function with variables subject to bounds. This algorithm uses gradient information; it is also called Newton Conjugate-Gradient. It differs from the *Newton-CG* method described above as it wraps a C implementation and allows each variable to be given upper and lower bounds.

Method *COBYLA* ([../optimize.minimize-cobyla.html#optimize-minimize-cobyla](#)) uses the Constrained Optimization BY Linear Approximation (COBYLA) method [R182], [10], [11]. The algorithm is based on linear approximations to the objective function and each constraint. The method wraps a FORTRAN implementation of the algorithm. The constraints functions 'fun' may return either a single number or an array or list of numbers.

Method *SLSQP* ([../optimize.minimize-slsqp.html#optimize-minimize-slsqp](#)) uses Sequential Least Squares Programming to minimize a function of several variables with any combination of bounds, equality and inequality constraints. The method wraps the SLSQP Optimization subroutine originally implemented by Dieter Kraft [12]. Note that the wrapper handles infinite values in bounds by converting them into large floating values.

### Custom minimizers

It may be useful to pass a custom minimization method, for example when using a frontend to this method such as `scipy.optimize.basinhopping` ([scipy.optimize.basinhopping.html#scipy.optimize.basinhopping](#)) or a different library. You can simply pass a callable as the `method` parameter.

The callable is called as `method(fun, x0, args, **kwargs, **options)` where `kwargs` corresponds to any other parameters passed to `minimize` (such as *callback*, *hess*, etc.), except the *options* dict, which has its contents also passed as *method* parameters pair by pair. Also, if *jac* has been passed as a bool type, *jac* and *fun* are mangled so that *fun* returns just the function values and *jac* is converted to a function returning the Jacobian. The method shall return an `OptimizeResult` object.

The provided *method* callable must be able to accept (and possibly ignore) arbitrary parameters; the set of parameters accepted by `minimize` may expand in future versions and then these parameters will be passed to the method. You can find an example in the `scipy.optimize` tutorial.

*New in version 0.11.0.*

## References

---

- [R174] (1, 2) Nelder, J A, and R Mead. 1965. A Simplex Method for Function Minimization. The Computer Journal 7: 308-13.
- [R175] (1, 2) Wright M H. 1996. Direct search methods: Once scorned, now respectable, in Numerical Analysis 1995: Proceedings of the 1995 Dundee Biennial Conference in Numerical Analysis (Eds. D F Griffiths and G A Watson). Addison Wesley Longman, Harlow, UK. 191-208.
- [R176] (1, 2) Powell, M J D. 1964. An efficient method for finding the minimum of a function of several variables without calculating derivatives. The Computer Journal 7: 155-162.
- [R177] (1, 2) Press W, S A Teukolsky, W T Vetterling and B P Flannery. Numerical Recipes (any edition), Cambridge University Press.
- [R178] (1, 2, 3, 4, 5, 6, 7, 8) Nocedal, J, and S J Wright. 2006. Numerical Optimization. Springer New York.
- [R179] (1, 2) Byrd, R H and P Lu and J. Nocedal. 1995. A Limited Memory Algorithm for Bound Constrained Optimization. SIAM Journal on Scientific and Statistical Computing 16 (5): 1190-1208.
- [R180] (1, 2) Zhu, C and R H Byrd and J Nocedal. 1997. L-BFGS-B: Algorithm 778: L-BFGS-B, FORTRAN routines for large scale bound constrained

optimization. ACM Transactions on Mathematical Software 23 (4): 550-560.

- [R181] (1, 2) Nash, S G. Newton-Type Minimization Via the Lanczos Method. 1984. SIAM Journal of Numerical Analysis 21: 770-778.
- [R182] (1, 2) Powell, M J D. A direct search optimization method that models the objective and constraint functions by linear interpolation. 1994. Advances in Optimization and Numerical Analysis, eds. S. Gomez and J-P Hennart, Kluwer Academic (Dordrecht), 51-67.
- [10] (1, 2) Powell M J D. Direct search algorithms for optimization calculations. 1998. Acta Numerica 7: 287-336.
- [11] (1, 2) Powell M J D. A view of algorithms for optimization without derivatives. 2007. Cambridge University Technical Report DAMTP 2007/NA03
- [12] (1, 2) Kraft, D. A software package for sequential quadratic programming. 1988. Tech. Rep. DFVLR-FB 88-28, DLR German Aerospace Center – Institute for Flight Mechanics, Koln, Germany.

## Examples

---

Let us consider the problem of minimizing the Rosenbrock function. This function (and its respective derivatives) is implemented in `rosen` (`scipy.optimize.rosen.html#scipy.optimize.rosen`) (resp. `rosen_der` (`scipy.optimize.rosen_der.html#scipy.optimize.rosen_der`), `rosen_hess` (`scipy.optimize.rosen_hess.html#scipy.optimize.rosen_hess`)) in the `scipy.optimize` (`./optimize.html#module-scipy.optimize`).

```
>>> from scipy.optimize import minimize, rosen, r
rosen_der
```

A simple application of the *Nelder-Mead* method is:

```
>>> x0 = [1.3, 0.7, 0.8, 1.9, 1.2]
>>> res = minimize(rosen, x0, method='Nelder-Mead', tol=1e-6)
>>> res.x
array([ 1.,  1.,  1.,  1.,  1.])
```

Now using the *BFGS* algorithm, using the first derivative and a few options:

```
>>> res = minimize(rosen, x0, method='BFGS', jac=
rosen_der,
...               options={'gtol': 1e-6, 'disp':
True})
Optimization terminated successfully.
        Current function value: 0.000000
        Iterations: 26
        Function evaluations: 31
        Gradient evaluations: 31
>>> res.x
array([ 1.,  1.,  1.,  1.,  1.])
>>> print(res.message)
Optimization terminated successfully.
>>> res.hess_inv
array([[ 0.00749589,  0.01255155,  0.02396251,  0
.04750988,  0.09495377], # may vary
       [ 0.01255155,  0.02510441,  0.04794055,  0
.09502834,  0.18996269],
       [ 0.02396251,  0.04794055,  0.09631614,  0
.19092151,  0.38165151],
       [ 0.04750988,  0.09502834,  0.19092151,  0
.38341252,  0.7664427 ],
       [ 0.09495377,  0.18996269,  0.38165151,  0
.7664427,   1.53713523]])
```

Next, consider a minimization problem with several constraints (namely Example 16.4 from [R178]). The objective function is:

```
>>> fun = lambda x: (x[0] - 1)**2 + (x[1] - 2.5)*
*2
```

There are three constraints defined as:

```
>>> cons = ({'type': 'ineq', 'fun': lambda x: x[
0] - 2 * x[1] + 2},
...        {'type': 'ineq', 'fun': lambda x: -x[
0] - 2 * x[1] + 6},
...        {'type': 'ineq', 'fun': lambda x: -x[
0] + 2 * x[1] + 2})
```

And variables must be positive, hence the following bounds:

```
>>> bnds = ((0, None), (0, None))
```

The optimization problem is solved using the SLSQP method as:



```
>>> res = minimize(fun, (2, 0), method='SLSQP', >>>
ounds=bnds,
...          constraints=cons)
```

It should converge to the theoretical solution (1.4,1.7).